CASE STUDY

Precision Pays

A Pre-Production A/B Experiment — Standardized Candidate Scoring – Product Area: Recruiting

July 2026 | Data Scientist: Francine Stephens

A UX test on candidate scoring surfaced a 23% accuracy gap – and a case for why more granular scoring design means more AI usage, not less.

Confidentiality Notice: *Specific numbers, user metrics, and internal terminology have been adjusted to protect Workday's proprietary information. The methods, reasoning, and outcomes represent the real project.*

Problem

Candidate evaluation is fragmented and uneven across Workday's ecosystem. Different teams rely on different scorecards, mental models, and data signals to assess candidates: some surfaces show letter grades, others show numeric percentages, and still others use trait-based scores. The one constant across all of them is inconsistent explainability. This fragmentation makes it harder for recruiters, PMOs, and hiring managers to compare applicants quickly and confidently, and it risks lowering the overall quality of talent selected.

Role & Context

I led the study design, statistical analysis, and reporting for this pre-production A/B experiment.

The study originated with a cross-functional team in the recruiting space — a PM, three UX designers, and a UX researcher — who had already conducted survey and interview research into how recruiters, PMOs, and hiring managers preferred to see candidate scores presented. That earlier research pointed toward numeric scoring as the stronger candidate, but the team needed evidence from real behavior, not just stated preference, before committing to a redesign.

I worked with the design technologist on the pre-production experimentation team to translate the two design variants the designers had produced into a testable experiment — defining the interaction logic needed to track UI behavior and measure time on task and success accurately.

Hypothesis & Design

Hypothesis: Replacing legacy letter grades with numeric percentage scores will help hiring managers, PMOs, and recruiters navigate candidate lists more easily, accurately, and quickly.

Method: 2,000 participants (hiring managers, PMOs, and recruiters), randomized to a single-exposure A/B test — Alpha (letter grade) vs. Numeric (percentage score) — reviewing a list of 8 candidates tiered A through D.

Task: Shortlist 3 candidates for an urgent contract Project Coordinator role that required Jira experience.

Success criteria: The most accurate shortlist selected both A-tier candidates plus the one B-tier candidate who had strong Jira experience.

Process & Decisions

Study design choice

The two B-tier candidates — Tyler (81%) and Anika (76%) — were deliberately chosen because they looked identical under a "B" letter grade but differed on the one skill that mattered for the role: Tyler was highly proficient in Jira, Anika had only basic experience. This pairing is what makes the mechanism visible later in the results — categorical ambiguity versus granular precision — rather than leaving the difference between conditions as a vague impression.

Metrics

Alongside the program's standard metrics — success rate, median time on task, and perceived ease of use, findability, confidence, and visual appeal — I added a post-task survey item asking participants directly how influential the score was to their shortlisting decision. I also defined a behavioral metric not captured by any single survey question: verification depth, calculated from how often participants interacted with UI elements that surfaced additional candidate detail — profile pages, AI summaries, resumes, and work history.

Proof-of-concept testing phase

Once the design technologist implemented the tracking logic, we launched a 10% sample as a proof-of-concept before scaling to the full study. This confirmed that interactive elements were correctly tagged in telemetry, pageviews were logging as expected, and the data pipeline was capturing behavior accurately before we committed the remaining sample.

Results

Say vs. do

Users said: Both groups self-reported near-identical reliance on the score — 79% of Alpha participants and 77% of Numeric participants rated it influential or extremely influential to their decision.

Users did: Actual accuracy diverged sharply — 82% success on Numeric, which was a significant 23% lift over the Alpha variant. Telemetry explains the gap: Numeric users viewed the top A-tier candidate's profile 2.4x more often (65% of participants) and ran more in-page searches for the Jira requirement (43% of participants on Numeric vs. 18% on Alpha).

Self-report data said the two variants were equivalent. Behavioral data said they weren't — the numeric variant was the one that actually drove a more thorough evaluation of the score and the surrounding candidate evidence.

Metric	Letter Grade (A)	Numeric Score (B)	Change (B vs. A)
Success rate	82%	67%	+23% lift

Metric	Letter Grade (A)	Numeric Score (B)	Change (B vs. A)
Median time on task	91 sec	118 sec	Not significant
Findability (% favorable)	77%	66%	−17% decline
Ease of use (% favorable)	84%	77%	—
Confidence (% favorable)	75%	80%	—
Self-reported score influence	79%	77%	— (equivalent)

Summary of core metrics, Letter Grade (n=990) vs. Numeric Score (n=1010).

Root cause: the categorical ambiguity trap

Letter grades collapsed real point differences — 81% vs. 76% — into an identical “B” badge. As a result, 15% of Alpha users shortlisted Anika despite her limited Jira experience, compared with only 6% in the Numeric condition. The grade was not just wrong — it just hid the one distinction the task actually depended on.

Unexpected trade-off: numeric scores drive thorough verification

Numeric scoring was not an unqualified win. Findability sentiment dropped 17 points (from 77% to 66% favorable) in the Numeric condition. The qualitative write-in feedback explains why: numeric precision prompted active verification — participants dug into profiles, resumes, and work history more often — rather than the passive acceptance a single letter badge invites. The friction was a byproduct of that added thoroughness, not a flaw in the interface itself.

Recommendations

- **Standardize on numeric (or hybrid) scoring** for this recruiting surface, given the accuracy gain it produced.
- **Redesign dashboard information density** to surface key evidence — skill matches, work history — inline on the candidate list, rather than requiring a click into each profile.
- **Increase visibility of the AI explainability feature:** only 19% of Alpha users and 13% of Numeric users opened it, despite it addressing exactly the kind of verification numeric users were doing manually.
- **Run a follow-up A/B test** once these changes are in place, to isolate how much of the findability friction the redesign with an inline or split-list view resolves.

Business Angle: Where Scoring UX Connects to Monetization

Beyond the UX outcome, this study surfaced a signal worth flagging for the business, separate from what the experiment was originally designed to measure. Numeric scores drove meaningfully deeper engagement with AI-generated content — profile and AI-summary views were 65% on Numeric vs. 27% on Alpha.

Under a consumption-based credits model, features that reliably pull users into more AI-assisted interactions — explainability panels, AI summaries, recommendation tabs — are not just a UX improvement; they generate usage. That reframes the low current discovery rate of the explainability feature (13–19%) as more than a UX gap — it is underused consumption sitting behind a discoverability problem.

Impact

The team responsible for this surface adopted the numeric scoring approach and is currently scoping a follow-up A/B test on explainability placement and easier discovery of candidate content. Based on these findings, PMs and designers invited me to present the results more broadly to the recruiting team members who own other candidate-scoring surfaces, to build the case for moving those surfaces toward numeric scoring as well.

Learnings

An unresolved trade-off

This surface has limited dashboard information density compared to other Workday surfaces. Other candidate-scoring surfaces already show more information on the main dashboard and may make explainability more discoverable by default — which could shrink, or shift, the gap seen here between letter and numeric scoring.

The value of pairing behavioral data with self-report

Self-report metrics alone would have missed the real story entirely. Both groups described their reliance on the score identically; only the behavioral data revealed that one group was actually doing the work the other believed it was also doing. It is a reminder that stated preference is useful, but it is not a substitute for evidence of what people actually do — self-report is subject to desirability bias and recall error in ways telemetry is not.